


FULLY OFFLINE | no cloud API | every model runs locally | 127.0.0.1 only

React + TypeScript Web UI

streaming | citations | corpus + document viewer

frontend/ built by Vite, served on :8000



FastAPI

/api/chat | /api/chat/stream | /api/documents/*

src/csrs/api/app.py

StreamLit UI

question | answer | doc list | model | reload


src/csrs/app.py on :8501

imports Pipeline directly

Pipeline facade - src/csrs/pipeline.py | index() ask() ask_stream()


neither interface imports Chroma, BM25 or Ollama

INGEST - Pipeline.index()



the index built above is what retrieval searches

QUERY - Pipeline.ask() / ask_stream()



Retrieval searches the rewritten query; generation always receives the user's original wording, so the answer addresses what was actually asked.

Ollama - http://127.0.0.1:11434

every orange stage above is a call to this, and to nothing else

Embedding (mandatory): nomic-embed-text | 274 MB | 768 dimensions

Answer models: llama3.2 2.0 GB (default) | qwen2.5:1.5b | gemma2:2b | phi4-mini | gemma4:e2b - all five selectable in both UIs

CSF-only index: 32 pages, 209 chunks | unchanged reload makes zero embedding calls | answers stream token by token over NDJSON

Add a standard by dropping a .pdf or .txt into docs/ and pressing Restart and Reload Documents - no code change, no restart

Proven, not asserted: the offline regression suite passes with Ollama pointed at a dead port - no module silently reaches the network.

Web fonts are vendored as local .woff2 files. There is no cloud LLM client anywhere in src/.

Legend User interface Backend / facade Local processing stage Ollama model call Index / retrieval Offline boundary Input