

```
{
  "alert_id": 23,
  "rule_id": "1:1122:16",
  "alert_message": "SERVER-WEBAPP /etc/passwd file access attempt",
  "snort_priority": 2,
  "classification": "Attempted Information Leak",
  "classtype": "attempted-recon",
  "query": "SERVER-WEBAPP /etc/passwd file access attempt Attempted Information Leak TCP http",
  "ranker_model": "openai/gpt-oss-120b",
  "model_rank": 4,
  "justification": "The alert indicates an HTTP request for /etc/passwd, which is classified as an attempted information leak and reflects simple reconnaissance rather than direct exploitation. The Snort priority is 2 (medium), but the activity is limited to information gathering with no immediate impact, justifying a low severity rating. Relevant fields such as alert_message, classification, priority, service, and packet_length support this assessment.",
  "anchored_rank": 3,
  "mismatch": false,
  "mismatch_explanation": null,
  "metrics_used": [
    "alert_message",
    "classification",
    "priority",
    "service",
    "packet_length"
  ],
  "sid_matching": {
    "snort_sid": 1122,
    "rag_matched_sid": 1122,
    "evidence_document": "snort_rule_1-1122.txt",
    "sid_match": true
  },
  "judge": {
    "model": "qwen/qwen3.6-27b",
    "score": 0.9,
    "reasoning": "The model assigned a rank of 4, which is within one step of the anchored ground truth of 3 (Snort priority 2 maps to rank 3). The justification is strong and evidence-based, correctly identifying the alert as reconnaissance rather than direct exploitation, which supports a slightly lower severity than the default mapping. The reasoning is complete and accurately reflects the rule documentation."
  }
}
```