

1. Run metadata

| Field                | Value                                                                   |
|----------------------|-------------------------------------------------------------------------|
| Provider             | Groq API                                                                |
| Model                | openai/gpt-oss-120b                                                     |
| Pipeline             | CSRS at /Users/rowdy/Projects/work/CIL/CSRS                             |
| Run id               | 20260813T093232Z                                                        |
| Sample               | /Users/rowdy/Projects/work/CIL/alert_sample_50.json (50 alerts)         |
| Snapshot             | /Users/rowdy/Projects/work/CIL/alert_rag_run.jsonl (50 rows)            |
| Options (all models) | temperature=0, max_completion_tokens=1200                               |
| Retrieval            | hybrid (dense top-20 + BM25 top-20, RRF k=60), rerank disabled, limit=8 |
| Embedding model      | nomic-embed-text (768-dim, task prefixes)                               |
| Corpus               | 4,039 documents / 6,518 chunks                                          |

6. Model versus Snort – confusion table (rows: Snort priority, columns: model rank)

| Snort \ model | -> 1 | -> 2 | -> 3 | -> 4 | -> 5 | Total |
|---------------|------|------|------|------|------|-------|
| 1             | 20   | 1    | 3    | 0    | 0    | 24    |
| 2             | 1    | 0    | 9    | 13   | 0    | 23    |
| 3             | 0    | 0    | 0    | 0    | 3    | 3     |

Exact match with Snort priority: 32/50 (64%)

Rank distribution: 1->21, 2->1, 3->12, 4->13, 5->3 vs sample composition 1->24, 2->23, 3->3